

Skip Lists: A Probabilistic Alternative to Balanced Trees

Presented by

Rishab Biswas - A125017
Siddharth Kumar - A125021

International Institute of Information Technology, Bhubaneswar

December 3, 2025



Outline

- 1 Linked Lists — The Starting Point
- 2 Skip Lists — The Solution
- 3 Simple Pseudocode
- 4 Variants of Skip Lists
- 5 Real-World Applications
- 6 Conclusion

Revisiting Singly Linked List

Search 67 → 5 steps



Why Linked Lists Are Not Enough

- Search = $O(n)$ → too slow for large data
- No random access
- Cache-unfriendly
- Poor concurrency support

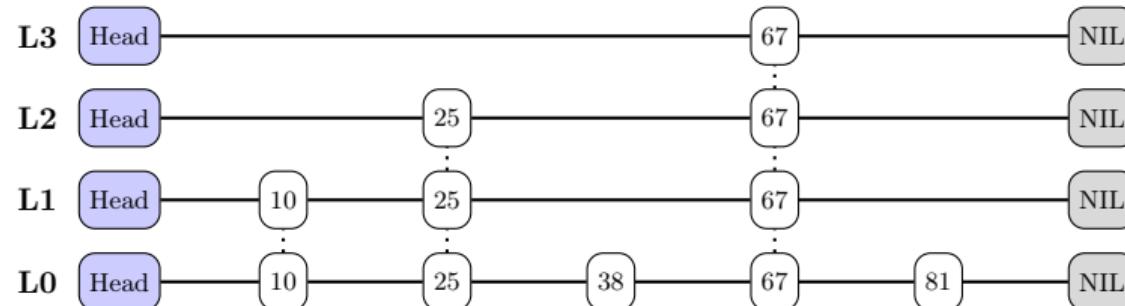
We need

Fast search + Simple implementation + Easy updates

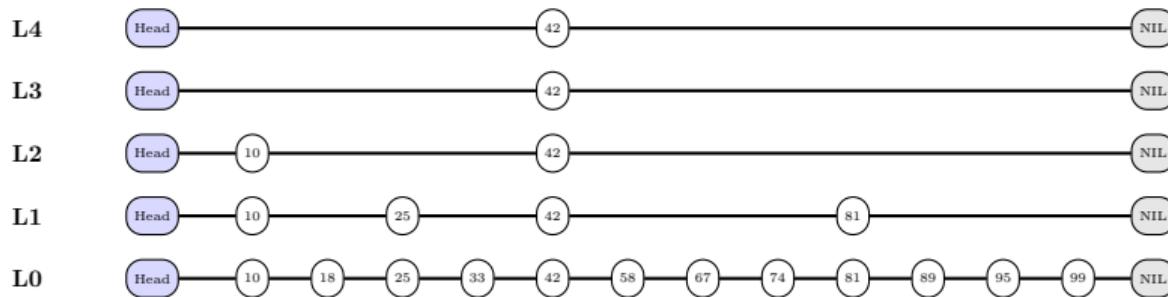
Skip List: Multi-Level Express Lanes

Brilliant Idea (William Pugh, 1990)

Randomly promote nodes to higher levels → create express lanes!

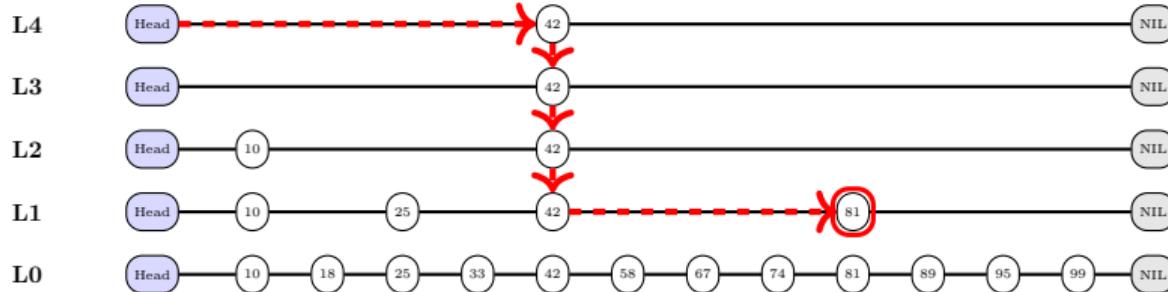


Realistic Example: Searching in a 12-Node Skip List



Slide 1: Full skip-list view before search.

Realistic Example: Search 81 in a 12-Node Skip List



Search 81 completed in 5 pointer moves (fast due to higher levels). Linked list needs 9 steps → Skip list wins!

Why Skip List Search is $O(\log n)$ (Easy Explanation)

Key Idea

Search moves **right** until the next value is too big, then **drops down**. There are $\log n$ levels, and only a constant number of right moves per level.

1. Expected Number of Levels

Nodes reach level i with probability p^i . Highest level $\approx \log_{1/p} n$. For $p = \frac{1}{2}$: $\log_2 n$ levels.

2. Expected Right Moves per Level

- Probability a node appears on level i is p .
- Looking backward, number of nodes scanned is geometric(p).

Expected right moves per level: $\frac{1}{p}$.

3. Total Expected Search Cost

Total expected steps:

$$\frac{1}{p} \log_{1/p} n + \frac{1}{1-p}$$

For $p = \frac{1}{2}$:

$$2 \log_2 n + 2$$

SEARCH(key)

- ① Start at top-left corner (highest level)
- ② While next node exists AND $\text{next.key} < \text{key}$ → move right
- ③ When stuck → drop down one level
- ④ Repeat until level 0
- ⑤ Return node if found

Insert — Random Height

INSERT(key, value)

- ① Perform SEARCH and remember one node per level (update[])
- ② If key exists → update value
- ③ Else:
 - Flip coin repeatedly → decide height
 - Create node with that height
 - Link it using update[] pointers

Randomness does the balancing automatically!

Delete — Just Fix Pointers

DELETE(key)

- ① Perform SEARCH and remember update[]
- ② If found → remove node from all its levels
- ③ Fix forward pointers (skip over it)

No rotations • No rebalancing • Super simple!

Space Complexity of Skip Lists ($p = 1/2$)

Expected number of nodes per level:

At level i , a node is promoted with probability p^i .

Expected nodes at level i = $n \cdot p^i$

For $p = \frac{1}{2}$:

$$n \left(\frac{1}{2}\right)^i$$

Total expected pointers in the entire structure:

$$\sum_{i=0}^{\infty} n \left(\frac{1}{2}\right)^i = n \cdot \sum_{i=0}^{\infty} \left(\frac{1}{2}\right)^i = n \cdot 2 = 2n$$

Final Result

Expected space complexity = $O(2n)$

Tradeoff: Time Complexity vs Space (Role of p)

Skip list performance depends on the promotion probability p .

Probability p	Height (Levels)	Space Used
High ($p \rightarrow 1$)	Tall structure	Many nodes per level (high space)
Low ($p \rightarrow 0$)	Short structure	Few levels (low space)

Search complexity:

$$O\left(\frac{1}{p} \log n\right)$$

Space complexity:

$$O\left(\frac{n}{1-p}\right)$$

Key Tradeoff

- Lower $p \rightarrow$ fewer levels \rightarrow **less space** but **slower searches**
- Higher $p \rightarrow$ more levels \rightarrow **faster searches** but **more space**
- $p = \frac{1}{2}$ is the sweet spot \rightarrow balanced time and space

Skip List Family

Deterministic → no randomness, guaranteed
 $O(\log n)$

Most Popular
Probabilistic + Concurrent

Concurrent/Lock-Free → Java's
ConcurrentSkipListMap

Finger Skip List → cached pointer (used in Redis)

Different p → $p = 1/e$ optimal, but $p = 0.5$ most common

Skip List vs Other Structures

Structure	Search	Insert/Delete	Code Complexity
Linked List	$O(n)$	$O(1)$ – $O(n)$	Very Simple
BST (unbalanced)	$O(n)$ worst	$O(n)$ worst	Simple
AVL / Red-Black	$O(\log n)$	$O(\log n)$	Complex
Skip List	$O(\log n)$ exp.	$O(\log n)$ exp.	Very Simple

Used in Production Everywhere

- **Redis** → Sorted Sets (ZSET)
- **Java** → ConcurrentSkipListMap/Set
- **LevelDB / RocksDB** → MemTable
- **Lucene / Solr** → Term dictionary
- **MongoDB WiredTiger, Cassandra, Intel TBB**

Why Industry Loves It Because

Simple + Fast + Excellent concurrency



Real-World Applications (Detailed)

1. Redis Sorted Sets (ZSET)

Used for:

- Leaderboards
- Time-series indexing
- Score-based range queries

Expected search time: $O(\log n)$

2. LevelDB / RocksDB MemTable

Skip lists provide:

- Fast in-memory writes
- Lock-free reads
- Efficient range scans

3. Java Concurrency

`ConcurrentSkipListMap/Set` offers:

- Non-blocking reads
- Sorted ordering
- Scalable multi-threaded performance

4. Search Engines (Lucene/Solr)

Used for:

- Term dictionary indexing
- Fast prefix/range lookups
- Real-time text search



Key Takeaways

Skip Lists = Best of Both Worlds

- Simplicity of Linked Lists
- Performance of Balanced Trees
- No complex rotations or rebalancing
- Expected $O(\log n)$ using only randomness

William Pugh's Legacy

"Replace complicated deterministic code with simple code + a little randomness"

References



Pugh, William.

Skip Lists: A Probabilistic Alternative to Balanced Trees.

Communications of the ACM, 33(6), 1990.



Herlihy, Maurice, and Shavit, Nir.

The Art of Multiprocessor Programming.

Morgan Kaufmann, 2008.



Redis Developers.

Sorted Sets (ZSET): Internal Implementation.

Available at: <https://redis.io/docs>



Google Developers.

LevelDB Architecture and MemTable Implementation.

Available at: <https://github.com/google/leveldb>



Apache Lucene.

Term Dictionary and Index Structures.

Thank You!

Thank You!

Questions?

Rishab Biswas Siddharth Kumar
IIIT Bhubaneswar December 3, 2025



IIIT Bhubaneswar
Imagine, Innovate, Inspire
(A university established by Government of Odisha)